

Interview Evaluation Guidelines

A structured evaluation framework for assessing candidates across all four GDS FDE levels. Designed to surface genuine FDE instinct — outcome orientation, technical depth, and the willingness to get close to consequence — not just credentials.

L1 · ANALYST

L2 · SENIOR

L3 · LEAD

L4 · MANAGER

DOCUMENT TYPE

Interviewer Reference ·
Confidential

DIMENSIONS PER LEVEL

5 Evaluation Dimensions

SIGNATURE QUESTIONS

3 per Level · FDE Instinct Probes

How to Use This Guide

This document is a live interviewing tool — not a checklist to complete after the fact. Work through each dimension during the interview, take notes in the provided space, and score immediately while memory is fresh. The three Signature FDE Questions at the end of each section are the most important part of any FDE interview.

FIVE DIMENSIONS PER LEVEL

Each level has five evaluation dimensions calibrated to what matters at that seniority. **Do not skip dimensions** — a strong technical score cannot compensate for a failing communication score in an FDE role. Use all five. Note evidence, not impressions.

SAMPLE QUESTIONS ARE STARTING POINTS

The questions listed are **prompts, not scripts**. Follow the energy of the conversation. The best FDE signal usually emerges from the follow-up: "And then what?" or "What would you do differently?" Push until you find the real answer, not the prepared one.

THE THREE SIGNATURE QUESTIONS

These are the most revealing questions in each interview. They are designed to be **impossible to fake with preparation alone**. They surface how a candidate thinks under pressure, handles failure, and whether they have the genuine FDE instinct — or are performing it.

SCORING & HIRE THRESHOLD

Each dimension is scored 1–4. A score of **3.0 average** across all dimensions is the minimum hire threshold. Any individual dimension score of **1 (Not Demonstrated)** is a hard flag that must be discussed in the debrief before proceeding. A score of 4 on a Signature Question carries double weight in the debrief discussion.

4

EXCEPTIONAL

Beyond level expectations. Shows genuine FDE instinct. Hire signal.

3

STRONG

Meets level requirements. Evidence is specific and credible. Proceed.

2

DEVELOPING

Partial evidence. Capability is present but not fully demonstrated. Probe further.

1

NOT DEMONSTRATED

No credible evidence. Red flag. Must be addressed in debrief.

The Sprint Runner

"We are not looking for someone who learns best in classrooms. We are looking for someone who shipped something before breakfast."

EXPERIENCE

0–2 Years

INTERVIEW DURATION

75 Minutes

PANEL SUGGESTED

L2 GDS FDE Senior + L3 or L4

CORE SIGNAL

Ships before seeking clarity

SUGGESTED INTERVIEW STRUCTURE

0–10 MIN

Warm-up + context setting

10–40 MIN

Technical + Delivery dimensions (D1, D2, D3)

40–60 MIN

Collaboration + FDE Spirit (D4, D5)

60–75 MIN

3 Signature Questions + candidate Q&A

01

Technical Foundations & Production Reality

Can they build something that works in the real world — not just a notebook? Is their code production-grade or academic?

1

2

3

4

SAMPLE QUESTIONS

Q1 Walk me through a project where you built something end-to-end — from idea to deployed. What did "deployed" mean in that context?

Listen for: Specificity. "Deployed to my laptop" is very different from "deployed to an API that others called."

Q2 You need to build a RAG system for 50,000 mixed-quality enterprise documents — inconsistent formats, multiple languages, no clear schema. What are your first three technical decisions and why?

Watch for: Do they ask the right clarifying questions? Do they think about chunking strategy, embedding quality, retrieval testing?

Q3 What is the difference between your code in a Jupyter notebook and your code in a production service? Show me the mental model, not just the answer.

Strong answer: Error handling, logging, versioning, schema validation, testability, reproducibility.

Q4 Walk me through the most complex API integration you've built. What made it hard and what did you do about it?

Red flag: Can't name a specific API or describes it in entirely abstract terms.

SCORING RUBRIC

4 Has specific, named production systems. Talks about failure modes, monitoring, schema drift. **Thinks in systems, not features.**

3 Has shipped something real. Understands the gap between demo and production. Some depth on technical tradeoffs.

2 Academic or tutorial-level experience. Understands concepts but hasn't shipped under real constraints. Vocabulary without evidence.

1 Cannot articulate what "production-grade" means. Experience is entirely theoretical. No real project to reference.

02

Delivery Instinct & Bias to Action

Do they ship before seeking perfection? Can they operate with incomplete information and still produce something useful?

1

2

3

4

SAMPLE QUESTIONS

- Q1 Tell me about a time you had incomplete requirements and had to ship anyway. What assumptions did you make? How did it turn out?

Strong answer: Named specific assumptions, documented them, shipped, then refined. Didn't wait.

- Q2 You're 6 hours into building something and you realise the approach fundamentally won't work. Walk me through what happens next.

Watch for: Paralysis vs. pivot. Do they escalate immediately or try to solve it first?

- Q3 Tell me about something you shipped that broke in production. What happened? What did you do in the first hour?

Red flag: Has never shipped anything that broke — means they've never shipped anything real.

- Q4 What does your GitHub commit history look like on a productive day? Walk me through a day you were proud of.

SCORING RUBRIC

- 4 Moves before the brief is perfect.**

Has specific stories of shipping imperfect things that worked well enough. Comfortable with iteration.

- 3** Has shipped under pressure. Understands the value of an imperfect working version over a perfect non-existent one.

- 2** Tends to seek clarity before acting. Not a strong signal of delivery instinct. May be cautious rather than action-oriented.

- 1** Waits for complete specs. Cannot describe a time they shipped under ambiguity. Comfort with uncertainty is absent.

NOTES

03

Problem-Solving Under Ambiguity

How do they approach problems they've never seen before? Do they reason from first principles or look for a template?

1

2

3

4

SAMPLE QUESTIONS

- Q1 A fraud operations team gets 31,000 alerts daily but only has capacity to review 8,000. Build them something in 6 hours. Talk me through your approach from first principles — no frameworks.

Listen for: How quickly do they decompose the problem? Do they ask about data availability, analyst workflow, what "better" looks like?

- Q2 You discover mid-sprint that the data you were told was clean has 40% missing values and inconsistent formats. Your delivery is in 4 hours. What do you do?

- Q3 What is the most interesting technical problem you have solved recently — interesting because of how hard it was

SCORING RUBRIC

- 4** Decomposes problems cleanly. **Asks the right clarifying questions before jumping in.** Shows creative lateral thinking alongside technical rigour.

- 3** Structured approach to new problems. Can reason without a template. Some evidence of first-principles thinking.

- 2** Tends to reach for familiar patterns. Struggles when a problem doesn't fit a known template. Gets stuck without a playbook.

- 1** Cannot structure an unfamiliar problem. No evidence of reasoning

to find the right solution, not just to implement it?

under genuine ambiguity.

Watch for: Feynman instinct. Do they enjoy the process of figuring things out? Or is it just task completion?

NOTES

04

Async Communication & Structured Handover

Can they communicate clearly in writing, without a synchronous session? Is their documentation a deliverable or an afterthought?

1

2

3

4

SAMPLE QUESTIONS

- Q1 You've built something overnight. Write me the Slack message you'd send to the onshore team at 7am — right now, in this interview — summarising what you built, what to test, and what's blocked. Go.

This is a live exercise. Watch for: clarity, structure, completeness. Does it tell the onshore team exactly what they need without a follow-up call?

- Q2 How do you handle a blocker that you discover at 2am when there's no one available from the onshore team to ask? Walk me through exactly what you do.

- Q3 What does your documentation look like for something you've built? Show me an example of README or technical note you'd be proud to hand over.

SCORING RUBRIC

- 4 Live Slack message is clear, structured, and complete. **No follow-up call needed.** Treats writing as a professional skill.
- 3 Good async communicator. Covers key points. Some room for tightening. The spirit of structured handover is present.
- 2 Tends to leave gaps. Async messages require follow-up questions to be useful. Documentation is light.
- 1 Cannot produce a clear written handover without prompting. Communication would create more work than it saves.

NOTES

05

FDE Spirit & Intrinsic Curiosity

Do they love figuring things out? Are they intrinsically motivated by the problem — or by the job title, the salary, and the process?

1

2

3

4

SAMPLE QUESTIONS

- Q1 What AI development or paper from the last three months has most changed how you think about something you're building? Walk me through the change.

Red flag: Nothing specific. "I follow the space generally." FDE people have a specific thing they're obsessed with.

- Q2 Tell me about something you built that nobody asked you to build. What problem were you solving and for whom?

- Q3 What is a problem in a domain outside of tech that you think AI could genuinely solve? How would you approach it?

Watch for: Commercial and domain curiosity alongside technical instinct. FDE requires both.

SCORING RUBRIC

- 4 Has specific things they are **obsessed with**. Lights up when talking about a problem. Has built things nobody asked for. Feynman instinct is unmistakable.
- 3 Genuine interest in the field. Has opinions. Follows developments. Shows curiosity beyond their job requirement.
- 2 Interest seems instrumental — motivated by career growth rather than genuine curiosity. Follows trends rather than forming views.

1 No intrinsic motivation evident.
Cannot articulate what excites them about AI or the problem space.

NOTES

3 Signature FDE Questions

LEVEL 01 · ANALYST · IMPOSSIBLE TO FAKE WITH PREPARATION ALONE

1 THE OVERNIGHT SPRINT SCENARIO

"It is 11:30pm. You receive a brief: 'Build a RAG system on this dataset — make it queryable — by 7am. The client ops team demos it at 8am.' You open the data and it is messier than expected — inconsistent formats, missing fields, mixed languages. There is no one to ask. Walk me through your next 7 hours — in detail."

WHAT YOU'RE LISTENING FOR

- How they handle the moment of discovering messy data — panic, pivot, or problem-solve?
- Do they scope down the problem immediately or try to solve for everything?
- Do they build a working "ugly" version first and clean it up — or stall seeking elegance?
- Do they document their assumptions or just push through?

GREEN SIGNALS

- ✓ Immediately scopes to the simplest thing that would demo well
- ✓ Builds a working pipeline on the cleanest subset first
- ✓ Documents assumptions as they go for the morning handover
- ✓ Shows calm, methodical urgency — not panic

RED SIGNALS

- ✗ First instinct is to message someone for clarity (no one is available)
- ✗ Wants to clean all the data before building anything
- ✗ Cannot describe the first 30 minutes with specificity
- ✗ The "demo" doesn't show up in their thinking until hour 6

2 THE HONEST FAILURE TEST

"Tell me about the last time you built something and then realised it was the wrong thing entirely — not broken, but wrong. How did you find out? What did you do? And what do you wish you had done differently?"

WHAT YOU'RE LISTENING FOR

- Can they be honest about failure without deflecting to external factors?
- Did they surface it themselves or wait to be told?
- Did they kill their own work or try to salvage it past the point of usefulness?

GREEN SIGNALS

- ✓ Has a specific, named story — not a hypothetical
- ✓ Surfaced the problem themselves before being asked
- ✓ Killed or pivoted their own work without being told to
- ✓ Reflection has a specific "what I'd do differently" — not a platitude

RED SIGNALS

- ✗ Cannot name a specific instance of genuine failure
- ✗ Frames every failure as someone else's misspecification
- ✗ Kept building something wrong because "that's what was asked"

→ Is the reflection specific and genuine — or generic and polished?		✗ "What I'd do differently" is vague or absent

3

THE FEYNMAN CURIOSITY PROBE

"What is a technical problem you are actively thinking about right now — outside of work or study? Not a project you were assigned. Something you cannot stop thinking about. Why does it interest you?"

WHAT YOU'RE LISTENING FOR	GREEN SIGNALS	RED SIGNALS
→ Is the curiosity genuine and specific — or performed and generic? → Do they have a point of view, or just an interest? → Does the problem reveal domain breadth beyond pure engineering? → Is this something they would stay up late to figure out?	✓ Specific problem with a specific reason they're stuck on it ✓ Can articulate what makes it hard — not just what makes it interesting ✓ Shows connection between the problem and their broader curiosity ✓ The energy in their answer is visibly different from the rest of the interview	✗ "I'm interested in AI generally" — no specific problem ✗ The problem is directly from their job description or coursework ✗ Curiosity is entirely instrumental — tied to career benefit ✗ Blank or long pause before manufacturing an answer

Overall Assessment Summary

DIMENSION	SCORE (1-4)	WEIGHT
D1 · Technical Foundations		×1.5
D2 · Delivery Instinct		×1.5
D3 · Problem-Solving		×1.0
D4 · Async Communication		×1.0
D5 · FDE Spirit		×1.0
Signature Questions (avg)		×2.0
Weighted Average Score		Min 3.0 to Hire

HIRE / NO-HIRE SIGNALS

— Strong Hire Signals

- ✓ Shipped something a stranger has used
- ✓ Has a specific, memorable failure story
- ✓ Overnight sprint answer is calm and methodical
- ✓ Curiosity probe reveals a genuine obsession

— Hard Pass Signals

- ✗ GitHub shows only tutorials and coursework
- ✗ Cannot describe production vs. notebook gap
- ✗ Paralysis on the overnight sprint question
- ✗ Curiosity is entirely instrumental

INTERVIEWER DECISION

HIRE

HOLD

PASS

The Technical Workstream Lead

"The shift from Analyst to Senior is not about knowing more. It is about owning more — and living with the consequences."

EXPERIENCE

3–5 Years

INTERVIEW DURATION

90 Minutes

PANEL SUGGESTED

L3 GDS FDE Lead + L4 Manager

CORE SIGNAL

Owns architecture + outcome simultaneously

SUGGESTED INTERVIEW STRUCTURE

0–10 MIN

Context + most recent production work

10–45 MIN

Architecture + Technical Judgment (D1, D2)

45–70 MIN

Business context + Mentoring + Patterns (D3, D4, D5)

70–90 MIN

3 Signature Questions + candidate Q&A

01

Advanced AI Architecture & System Design

Can they design end-to-end agentic systems — not just implement components? Do their architectural decisions reflect real production experience?

1

2

3

4

SAMPLE QUESTIONS

- Q1 Design a multi-agent triage system for a fraud operations team receiving 31,000 alerts per day with 8,000 analyst capacity. Walk me through the full architecture — agents, tools, state management, guardrails, and how you'd measure success.

Listen for: Do they ask about false positive tolerance, analyst workflow, regulatory constraints? Do they design for failure modes from the start?

- Q2 What are the three failure modes of RAG that most production teams discover too late? For each, how do you architect your way out of it from Day 1?

- Q3 When would you explicitly choose NOT to use an agentic architecture for a client problem? What are the signals that tell you a simpler system is right?

Strong answer: Knows when agentic is over-engineered. Most weak candidates always reach for agents.

- Q4 You need to build an explainability layer for an LLM decision system that a bank's compliance team must

SCORING RUBRIC

- 4 Designs for failure modes before designing for features. **Architecture decisions are explained with "because" not just "how."** Shows awareness of cost, latency, and governance alongside capability.

- 3 Solid system design. Covers the key components. Some depth on tradeoffs. Production-aware even if not exhaustive.

- 2 Covers the components but not the connections or failure modes. Design works on paper but would surprise them in production.

- 1 Architecture is shallow or generic. Would not survive production constraints. Confuses knowing the tools with knowing system design.

audit. What does that look like technically and what are the non-negotiables?

NOTES

02

Technical Judgment & Workstream Ownership

Have they made architecture decisions with real consequences? Can they own a broken system and fix it — without being told how?

1

2

3

4

SAMPLE QUESTIONS

- Q1 Tell me about a technical decision you made — not contributed to, but owned. What was the risk, how did you make the call, and what happened?

Listen carefully for "I" vs "we." FDE Senior owns decisions, not just contributes to them.

- Q2 You inherit a "working" codebase. It demos beautifully. You suspect it will not survive production load or edge cases. You have 2 weeks to stabilise it before a client demo. What is your process — step by step?

- Q3 Describe a system you built that worked in the demo and failed in production. What did you miss? How did you find out? What was fixed and what was a fundamental rethink?

Red flag: No production failures = no production experience.

SCORING RUBRIC

- 4 Has specific stories of **owning the hard call**. Knows the difference between being responsible and being accountable. Production failure story is detailed and honest.
- 3 Has owned meaningful technical scope. Can describe the decisions and consequences. Honest about limitations.
- 2 Contributed to decisions but rarely owned them. Scope of ownership has always been bounded and supervised.
- 1 Has always worked under close direction. Cannot describe a moment of independent technical ownership.

NOTES

03

Business Context Awareness & Translation

Can they connect what they build to why the client cares? Can they explain a vector database to a VP of Operations without condescension?

1

2

3

4

SAMPLE QUESTIONS

- Q1 How do you know if the AI system you built actually changed a business outcome? What is the difference between a technical metric and a business metric — and how do you make sure you're tracking the right one?

- Q2 Tell me about a time you pushed back on a requirement because you believed it was technically unsound or would not achieve the business goal. What happened?

Strong signal: They pushed back with evidence, not opinion — and they were right. Or they were wrong and learned from it.

- Q3 Write me — right now — a 3-sentence summary of a system you built, intended for a VP of Operations who has no technical background. Go.

SCORING RUBRIC

- 4 **Naturally leads with "so what" before "how."** VP summary is clear and outcome-focused without being dumbed down. Business and technical lenses switch fluidly.
- 3 Understands why business outcomes matter. Can translate with effort. VP summary is decent but leans slightly technical.
- 2 Business lens is secondary — technical framing dominates. VP summary is hard to follow or uses unexplained jargon.

Live exercise. Watch for: Does it lead with what it does for them, or how it works?

- 1 Cannot write a non-technical summary without producing a technical one. "They don't need to understand how it works" is their position.

NOTES

04

Mentoring & Team Elevation

Have they made someone else better? Can they describe the specific mechanism — not just that they "helped" a junior?

1

2

3

4

SAMPLE QUESTIONS

- Q1 Tell me about a junior engineer you developed. Not just helped — developed. What specifically changed because of what you did? What was the before and after?

Listen for specificity. "I reviewed their code" is not development. "I put them on a problem slightly above their level and stayed visible for when they got stuck" is development.

- Q2 Describe a code review you gave that you are proud of. What was the feedback, how did you deliver it, and what changed in their work afterwards?

- Q3 A junior on your team delivers something that works but is architecturally bad. The deadline was yesterday. What do you do?

Watch for: Ship now, fix later mentality vs. "let me fix it" vs. using it as a teaching moment — in what order?

SCORING RUBRIC

- 4 Has a named person who grew under their influence. Can describe **the specific mechanism** — what they did, not just that they did it. Development is a natural instinct, not a job requirement.

- 3 Has mentored effectively. Genuinely interested in others' growth. Code reviews are a teaching tool, not just a quality gate.

- 2 Has helped juniors but through proximity, not intent. Mentoring is reactive — triggered by someone asking, not proactively offered.

- 1 Has not developed anyone. Views junior work as something to be corrected, not developed. No mentoring instinct visible.

NOTES

05

FDE Pattern Recognition & Field Instinct

Do they recognise the common patterns of failure in enterprise AI before being told about them? Have they accumulated real field wisdom?

1

2

3

4

SAMPLE QUESTIONS

- Q1 What are the patterns you reliably see in enterprise AI projects that fail — not at the technology level, but at the people and process level? What is the first sign?

Strong answer references specific patterns: overscoped POCs, no eval framework, no owner post-launch, model team isolation, demo-to-production gap. Not textbook — lived.

- Q2 Tell me about something you built that you would never build the same way again. What changed — in you, not in the technology?

SCORING RUBRIC

- 4 Has **accumulated distinct patterns from real engagements**. Can predict what will go wrong from early signals. Field wisdom is distinct from textbook knowledge — it has texture and specificity.

- 3 Has clear opinions formed through real experience. Pattern recognition is present, if not yet fully calibrated.

Q3 What is the single most underestimated challenge in deploying agentic AI in a regulated enterprise environment? Not a technical answer — a human and organisational one.

The best candidates name something specific: model risk governance, change management with ops teams, or the feedback loop from operations to model development.

2 Opinions are largely borrowed from articles and podcasts rather than lived experience. Patterns are generic rather than specific.

1 No field wisdom. Answers to these questions would be found in any AI deployment blog post. No lived accumulation visible.

NOTES

3 Signature FDE Questions

LEVEL 02 • SENIOR • DESIGNED TO SURFACE OWNERSHIP INSTINCT

1

THE PRODUCTION HONESTY TEST

"You have built a system that demos beautifully. Stakeholders are happy. The client is pleased. You run the first real production test with live data and discover a 34% hallucination rate on a critical output. Nobody else has noticed yet. What do you do in the next 24 hours?"

WHAT YOU'RE LISTENING FOR

- How fast do they surface it — and to whom?
- Do they investigate first or immediately escalate?
- Do they own it or soften it ("the model has some limitations...")?
- What is their communication strategy — who gets told what and in what order?

GREEN SIGNALS

- ✓ Documents the finding thoroughly before raising it
- ✓ Surfaces it honestly and promptly — no softening
- ✓ Has a proposed fix or investigation path ready when they raise it
- ✓ Does not wait for someone else to notice

RED SIGNALS

- ✗ First instinct is to keep building and "improve it later"
- ✗ Frames it as "the model has limitations" rather than "my system has a problem"
- ✗ Escalates without first investigating
- ✗ "I'd check with the team" without specifying what they'd check

2

THE INFLUENCE WITHOUT AUTHORITY TEST

"Describe a time you were technically right but couldn't get anyone to listen. What happened, what did you do, and — looking back — what would you do differently?"

WHAT YOU'RE LISTENING FOR

- Is the story real and specific — or a generic answer about "communicating better"?
- Do they reflect on their own role in the failure to influence, or only blame others?

GREEN SIGNALS

- ✓ Specific, named incident with real stakes
- ✓ Self-aware about what they could have done differently
- ✓ Shows that they tried multiple approaches before concluding it failed

RED SIGNALS

- ✗ Cannot name a specific incident
- ✗ Story positions them as purely right and others as purely wrong
- ✗ No reflection on their own communication or approach

→ What was the actual mechanism they used to try to change minds? → Did they give up, escalate, or find another way in?	✓ "What I'd do differently" is a specific mechanism, not a platitude	✗ Resolution was "I got overruled and moved on" with no further action

3

THE MID-ENGAGEMENT PIVOT TEST

"You are three weeks into an engagement. The onshore FDE Lead sends you a message at 9pm: 'I think we might be building the wrong thing entirely.' You have two workstreams running, one GDS Analyst on your team, and a client demo in five days. What is your response — and what do you actually build next?"

WHAT YOU'RE LISTENING FOR	GREEN SIGNALS	RED SIGNALS
→ How fast do they orient around the new information? → Do they protect their existing work or immediately interrogate it? → How do they manage the GDS Analyst in this moment of uncertainty? → What is their first question back to the FDE Lead?	✓ First response is a clarifying question, not a reassurance ✓ Immediately assesses what of the existing work is salvageable vs. sunk cost ✓ Thinks about the Analyst's clarity and task before going to sleep ✓ Has a working hypothesis about what to build instead within the hour	✗ First instinct is to defend what was already built ✗ "We can't pivot this late" is their framing ✗ Waits for the morning call before doing anything ✗ Does not mention the Analyst — only thinks about their own workstream

Overall Assessment Summary

DIMENSION	SCORE (1-4)	WEIGHT	HIRE / NO-HIRE SIGNALS
D1 · AI Architecture & System Design		×1.5	— Strong Hire Signals ✓ Owns a production system others depend on ✓ Surfaces bad news before being asked ✓ Can name someone who grew under them ✓ Pattern recognition is field-earned, not textbook — Hard Pass Signals ✗ "We" for every ownership question ✗ Protects broken work past the point of usefulness ✗ Business translation is jargon-heavy or absent
D2 · Technical Judgment & Ownership		×1.5	
D3 · Business Context & Translation		×1.0	
D4 · Mentoring & Team Elevation		×1.0	
D5 · FDE Pattern Recognition		×1.0	

Signature Questions (avg)		×2.0
Weighted Average Score		Min 3.0 to Hire

✗ No real experience mentoring anyone

INTERVIEWER DECISION

HIRE

HOLD

PASS

The POD Technical Anchor

"At this level the question is not: can they build it? It is: can they make every other engineering decision on the engagement better — including ones they didn't make?"

EXPERIENCE

5–8 Years

INTERVIEW DURATION

100 Minutes

PANEL SUGGESTED

L4 GDS Manager + Onshore Director/SM

CORE SIGNAL

Technical authority without being in the room

SUGGESTED INTERVIEW STRUCTURE

0–10 MIN

Context + most complex system designed

10–50 MIN

Architecture + Governance + Client Influence (D1, D2, D3)

50–75 MIN

Team leadership + Commercial awareness (D4, D5)

75–100 MIN

3 Signature Questions + candidate Q&A

01

Agentic Architecture at Enterprise Scale

Can they design for the full enterprise context — governance, legacy integration, reliability, compliance — not just the AI components?

1

2

3

4

SAMPLE QUESTIONS

- Q1 Design the full technical architecture for an agentic fraud alert triage system inside a regulated bank: from alert ingestion through MRM-compliant AI decisioning to analyst workflow integration. What are your five most critical design decisions and why?
- Specifically watch for: audit trail design, human-in-the-loop placement, hard guardrail architecture, and how they handle the MRM classification question.*
- Q2 A bank's model risk team has classified your agentic post-processing layer as a "shadow model" requiring 6–8 months of SR 11-7 validation. You need it live in 8 weeks. Walk me through your redesign strategy — specifically what you change technically and how you reframe it for MRM approval.
- Q3 What does production-grade look like for an agentic system that a bank's compliance team will audit quarterly? Go beyond "logging and monitoring" — what are the governance design patterns?

SCORING RUBRIC

- 4 Designs governance into the architecture from the start. **Knows how to reframe an AI system for a different regulatory track without sacrificing the technical design.** Audit trail and explainability are first-class design elements.
- 3 Strong architecture. Governance awareness is present and practical. Knows what MRM cares about even if not in exact detail.
- 2 Good AI architecture but governance is bolted on rather than designed in. Compliance questions reveal gaps.
- 1 Architecture would not survive a compliance review. Governance is an afterthought. MRM question reveals no practical understanding.

02

Governance & Compliance Navigation

Have they navigated a real regulatory wall — not by fighting it but by redesigning around it? Can they speak MRM's language?

1

2

3

4

SAMPLE QUESTIONS

- Q1 Walk me through a specific time you encountered a regulatory or compliance blocker in an AI deployment. What did you change — technically and in how you framed it — to get approval? What was the exact mechanism, not the headline outcome?
- Q2 How do you design an AI system that satisfies a CISO, a Chief Risk Officer, and a Chief Data Officer simultaneously — when their concerns are genuinely in conflict? Walk me through a specific example.
- Q3 What is the most important thing to do in the first 48 hours when you discover that a compliance team has classified your AI system as requiring a governance track that will take 6 months?

Strong answer: Investigate the classification criteria before escalating. Then redesign before negotiating. Then bring data to the approval meeting, not arguments.

SCORING RUBRIC

- 4 Has a specific, lived governance navigation story.** Redesigned the system to earn approval — didn't just argue for a waiver. Knows what regulators actually care about vs. what they say they care about.
- 3** Has navigated compliance constraints. Understanding is practical if not deep. Would find the right path with some additional support.
- 2** Has seen governance constraints but worked around them rather than through them. Would need senior support in a live MRM conversation.
- 1** No practical governance experience. Frames compliance as a blocker to navigate around rather than a design constraint to design through.

03

Client Influence from Distance

Can they hold a client technical room from a GDS position? Have they changed a client's technical direction without being physically present?

1

2

3

4

SAMPLE QUESTIONS

- Q1 Tell me about a time you changed a client's or senior stakeholder's technical direction from a remote or offshore position. What was the mechanism — how did you establish credibility without being in the room?
- Q2 The client's architecture team has proposed an approach that you know will fail under production load. They are not your direct reports and they have more seniority. You are offshore. What is your strategy over the next 5 days?
- Q3 What does building trust with a sceptical client engineering team look like from a GDS position? Walk me through the first 4 weeks — specifically what you do and what changes.

SCORING RUBRIC

- 4** Has specific stories of **influencing client direction from offshore.** Understands that influence from distance requires more evidence and less opinion than influence in the room. Uses work quality as the primary credibility mechanism.
- 3** Has influenced technical decisions with senior stakeholders. Remote influence is understood as a specific skill requiring deliberate approach.
- 2** Has influenced decisions when physically present. Remote influence strategy is underdeveloped. Would need onshore support for high-friction client conversations.

- 1 Has not influenced client technical decisions at all. Views this as "the onshore team's job." Not ready for an L3 FDE role.

NOTES

04

Team Leadership & Delivery Quality Standard

Do they set the standard by example — or by process? Can they describe the mechanism through which they raise a team's output quality?

1

2

3

4

SAMPLE QUESTIONS

- Q1 How do you set the technical quality standard for a team you lead — not in terms of process documents, but in terms of what you personally do? Walk me through a concrete example.
- Q2 Your GDS delivery is visibly in trouble — two weeks from a milestone, output quality has dropped, the team is demoralized, and the onshore Lead is asking questions. What is your diagnosis process and your first 72-hour plan?
- Q3 How do you know when an engineer on your team is ready to take on more than they are currently handling? What specific signals do you look for — and what do you do first when you see them?

SCORING RUBRIC

- 4 Sets standards through visible example — **the team sees what "great" looks like in their Lead's own work.** Diagnoses team problems through observation and data, not assumptions. Developing others is active and deliberate.
- 3 Clear leadership instincts. Sets standards deliberately. Has navigated a team performance challenge successfully.
- 2 Leadership is through process rather than example. Has led teams but quality standard is enforced through review gates, not culture.
- 1 Has not led a team through a delivery challenge. Quality standard is undefined or delegated entirely upward.

NOTES

05

Commercial & Engagement Awareness

Do they understand that their work has a commercial context? Can they see the connection between GDS delivery and client value creation?

1

2

3

4

SAMPLE QUESTIONS

- Q1 How do you know if a GDS FDE engagement is delivering commercial value to the client? What metrics do you track beyond technical delivery milestones?
- Q2 What does a good exit look like for a GDS FDE engagement? How do you specifically prepare the client team to own the system after you leave — and how do you know when they are ready?

SCORING RUBRIC

- 4 **Naturally frames their work in terms of client value, not task completion.** Exit planning is a first-class concern from Day 1. Scope conversations are handled with commercial maturity.
- 3 Commercial awareness is present and growing. Connects work to client outcomes. Exit planning is understood even if not yet second nature.

03 You are 3 weeks into an engagement and you realise the scope was scoped wrong — the client has a significantly larger problem than the SOW covers and cannot be fixed properly within it. What do you do and who do you tell?

2 Commercially unaware. Measures success in technical terms only. Exit is someone else's problem.

1 No commercial awareness visible. Has never thought about their work in terms of client value or engagement economics.

NOTES

3 Signature FDE Questions

LEVEL 03 · LEAD · DESIGNED TO SURFACE TECHNICAL COURAGE AND HUMAN INTELLIGENCE

1

THE HOSTILE EXPERT TEST

"The model team at your client has worked on their fraud model for four years. They are brilliant, protective, and convinced that the problems the FDE team is finding are operational issues — not model failures. You have field data that strongly suggests otherwise. How do you have this conversation — and how do you make them want to hear it?"

WHAT YOU'RE LISTENING FOR

- Do they frame this as a confrontation or a collaboration?
- What is the mechanism — how do they actually change minds?
- Do they understand what the model team is afraid of losing?
- Is their approach data-first — or argument-first?

GREEN SIGNALS

- ✓ Leads with evidence, not conclusions — shows before tells
- ✓ Understands the model team's identity is in their work — addresses that explicitly
- ✓ Frames their field data as enrichment of model team knowledge, not critique
- ✓ Patient — willing to be heard over weeks, not forced in one meeting

RED SIGNALS

- ✗ Escalates to leadership as the primary strategy
- ✗ Frames it as "getting them to accept" the data
- ✗ Does not consider what the model team is protecting
- ✗ "I'd just show them the numbers" without a human strategy

2

THE ARCHITECTURE INTEGRITY TEST

"You are the most senior technical person on a GDS engagement. Six weeks in, you discover that the architecture the client signed off before your engagement began will fail catastrophically under production load in a way that causes real business harm. The onshore Partner is not technical. The client CTO is defensive. You have no authority here. Walk me through exactly what you do."

WHAT YOU'RE LISTENING FOR

- Do they have the courage to surface it despite the political

GREEN SIGNALS

- ✓ Builds the evidence case thoroughly before raising it

RED SIGNALS

- ✗ Waits for it to become an incident before raising it

cost?	✓ Brings a proposed solution, not just the problem ✓ Finds an indirect route to the CTO through a trusted relationship ✓ Their primary framing is "client protection" — not "I was right"	✗ Raises it without a proposed alternative ✗ Uses the CTO's defensiveness as a reason not to act ✗ Defers entirely to the onshore Partner to handle it
→ What is their brief-building strategy before they raise it? → How do they handle the defensive CTO without damaging the relationship? → Is their primary concern the client's outcome or their own position?		

3

THE SELF-AWARENESS PROBE

"What is the hardest thing about leading highly technical work from a GDS position that you have genuinely not figured out yet? Not a polished answer — a real one."

WHAT YOU'RE LISTENING FOR	GREEN SIGNALS	RED SIGNALS
→ Is the answer specific and honest — or a disguised strength? → Does the challenge they name reveal authentic self-knowledge? → Are they curious about solving the gap — or just naming it? → Does their answer tell you something true about them that nothing else in the interview did?	✓ Specific, named challenge that feels genuinely vulnerable ✓ Has already tried some approaches and can describe what hasn't worked ✓ Shows that they think about their own development as deliberately as they think about their team's ✓ The answer reveals something real — you learn something from it	✗ "I'm still working on delegating more" — a generic non-answer ✗ Pivots to a strength dressed as a weakness ✗ Long pause followed by a clearly manufactured answer ✗ Cannot name anything genuinely unsolved — signals lack of self-reflection

Overall Assessment Summary

DIMENSION	SCORE (1–4)	WEIGHT	HIRE / NO-HIRE SIGNALS
D1 · Agentic Architecture at Scale		×1.5	— Strong Hire Signals ✓ Surfaces architecture problems before they become incidents ✓ Has a real, lived governance navigation story ✓ Changed a client technical decision from a GDS position ✓ Self-awareness probe answer is specific and vulnerable — Hard Pass Signals ✗ Governance is always someone else's problem
D2 · Governance & Compliance Navigation		×1.5	
D3 · Client Influence from Distance		×1.0	
D4 · Team Leadership & Quality Standard		×1.0	

D5 · Commercial & Engagement Awareness		×1.0
Signature Questions (avg)		×2.0
Weighted Average Score		Min 3.0 to Hire

- ✗ Cannot describe influencing a senior without authority
- ✗ Architecture does not survive compliance questioning
- ✗ Self-awareness probe produces a generic non-answer

INTERVIEWER DECISION

HIRE

HOLD

PASS

The GDS Mission Commander

"The question at this level is not: are they impressive? It is: do they make everyone around them more capable — including people they will never meet?"

EXPERIENCE

8–12 Years

INTERVIEW DURATION

2 Sessions · 90 min each

PANEL SUGGESTED

Onshore Partner + GDS Director + L3 Lead

CORE SIGNAL

Multiplies capability without losing credibility

SUGGESTED TWO-SESSION STRUCTURE

SESSION 1 · 0–90 MIN

**Technical authority
+ People leadership
(D1, D2)**

SESSION 1 · 60–90
MIN

**Practice building +
Commercial (D3, D4)**

SESSION 2 · 0–60 MIN

**FDE culture +
Engagement
complexity (D5) +
Signature Qs**

SESSION 2 · 60–90
MIN

**Panel debrief +
Candidate Q&A with
full panel**

01

Technical Authority & Currency

Are they still technically credible enough to lead engineers? Do they maintain hands-on contact with the work — or have they drifted into pure management?

1

2

3

4

SAMPLE QUESTIONS

Q1 What is the most significant development in agentic AI in the last 6 months? How have you incorporated it — or challenged it — in your practice or an active engagement?

Red flag: Vague or behind the current state. Strong signal: Has a specific view on a specific development with a specific implication.

Q2 Walk me through your last hands-on technical contribution to a live engagement. What were you building? What decisions did you make — and how did you know they were yours to make?

Q3 How do you maintain technical currency when your primary job is leading people? What is your actual practice — specifically, not aspirationally?

Q4 Review this architecture diagram with me (provide a live architecture challenge). What are the three most important issues you see? What questions would you ask before proceeding?

SCORING RUBRIC

4 Technically current. Has hands-on time built into their working week. **Could step into a Lead role on an engagement and be effective within 48 hours.** Technical credibility is earned weekly, not borrowed from past experience.

3 Technically credible. Some drift from the coalface but maintains enough contact to lead and review effectively. Would benefit from structured hands-on time.

2 Technical credibility is waning. Last substantive hands-on contribution was more than 12 months ago. Would struggle to review complex architecture without support.

1 Has left the technical work behind entirely. Would not be credible in a technical review with senior client

Live technical exercise. At this level, they must be able to find issues — not just have an opinion on them.

engineers. Cannot lead what they can no longer do.

NOTES

People Leadership & Talent Development

02

Have they produced other great engineers — deliberately, not by accident? Can they describe the specific mechanisms they use to develop people?

1

2

3

4

SAMPLE QUESTIONS

- Q1 Tell me about an engineer who is now better than you were at their level. What did you do — specifically — that made the difference? What was hard about developing them?

The best answer names a person, a specific challenge in their development, a specific mechanism used, and an honest reflection on cost or difficulty.

- Q2 Tell me about an engineer you had to manage out. What happened, how did you know, what did you do, and what would you do differently?

Anyone who has led teams has this story. If they don't — they haven't led teams, or they avoided the hard conversation.

- Q3 What does your calibration process look like when assessing whether an L2 GDS FDE Senior is ready for an L3 Lead role? What specific evidence do you look for — not potential, but demonstrated evidence?

- Q4 How do you develop someone who is technically exceptional but struggles to communicate their work in business terms — and is resistant to feedback about it?

SCORING RUBRIC

- 4 Has named people whose careers were shaped by their leadership. **Development is a deliberate practice with specific mechanisms** — not good luck or proximity. Has also managed out effectively and learned from it.

- 3 Has developed people meaningfully. Calibration process is clear and evidence-based. Has navigated a difficult performance conversation.

- 2 Has managed teams but development is informal and inconsistent. Calibration lacks rigour. Has avoided hard performance conversations.

- 1 Has not developed anyone intentionally. Cannot describe the mechanism. Has never managed out — or refuses to acknowledge it. Not ready for this level.

NOTES

Practice Building & Commercial Leadership

03

Have they built something that outlasts their tenure? Do they understand the commercial context of GDS FDE deeply enough to own it?

1

2

3

4

SAMPLE QUESTIONS

- Q1 What reusable asset, pattern, methodology or accelerator did you build that benefited teams and engagements beyond your own immediate team? What was the mechanism for adoption, and what is its current state?

- Q2 Walk me through how you would scope GDS FDE involvement in a new engagement proposal — from your first read of the client brief to the final team plan. What

SCORING RUBRIC

- 4 **Has built something that teams are still using.** Proposal process is detailed, evidence-based, and includes specific pushback examples. Commercial challenges are faced, not avoided.

- 3 Has contributed meaningfully to practice development. Commercial

factors do you weight? What do you typically push back on?

Q3 Tell me about an engagement that was commercially challenging — where GDS delivery costs and client value were out of alignment. How did you diagnose it and what did you do?

awareness is solid if not sophisticated. Scoping process is clear.

2 Practice development has been incremental and local. Commercial conversations are delegated upward. Scoping is reactive, not proactive.

1 No practice-building visible. Commercial context is absent from their thinking. Would need significant support to own a GDS FDE commercial relationship.

NOTES

High-Stakes Engagement Navigation

04

Have they navigated a truly hard engagement — one with political, commercial, and technical failure modes simultaneously? What did they learn?

1

2

3

4

SAMPLE QUESTIONS

Q1 Tell me about the most difficult engagement you have led. What specifically made it hard — not technically, but organisationally and politically. What did you do that you had never done before? What would you do differently?

Q2 You are 6 weeks into a flagship EY engagement. The client CTO has lost confidence in EY's delivery. The onshore Partner asks you — specifically — to take the technical lead on recovery. Walk me through your first 72 hours and your 2-week plan.

Q3 When is the right moment to tell a client that the problem they hired you to solve is not actually their real problem? How do you have that conversation without losing the engagement?

SCORING RUBRIC

4 Has a specific hard engagement with real stakes. **Recovery plan is detailed, sequenced, and draws on hard-won pattern recognition.** "What I'd do differently" reflects genuine learning, not performance.

3 Has navigated difficult engagements. Recovery instincts are sound. Real problem conversation is handled with maturity.

2 Has experienced difficult engagements but not led through them. Complexity was handled by others or escalated upward. Recovery plan is generic.

1 No high-stakes engagement experience. Cannot describe a scenario where their leadership was the decisive factor in a difficult outcome.

NOTES

FDE Culture & the Multiplier Effect

05

Do they understand the FDE culture well enough to protect it at scale? Can they distinguish genuine FDE instinct from its performance?

1

2

3

4

SAMPLE QUESTIONS

SCORING RUBRIC

4 Can distinguish FDE culture from FDE process. **Knows that the instinct is**

Q1 What is the most important thing you do to prevent a GDS FDE team from sliding back into a traditional delivery centre mentality — especially under commercial and schedule pressure?	preserved through example and consequence, not enforcement. Has a specific mechanism for detecting cultural drift before it becomes visible in delivery.
Q2 How would you assess whether someone new to the GDS FDE team has genuine FDE instinct versus is performing FDE because it is required of them? What are the signals — in how they work, not what they say?	3 Clear view of what FDE culture means and how to protect it. Some sophistication about how it erodes under pressure.
Q3 You have 40 engineers across 6 engagements. Three months in, you notice the teams are technically excellent but gradually losing the "question whether we're building the right thing" instinct — nobody has done anything wrong, they're just executing beautifully toward the wrong outcome. How do you diagnose and address this?	2 Values FDE culture in principle but would address drift through process changes rather than cultural leadership. Mechanism is compliance, not conviction.
	1 Cannot articulate what makes FDE culture distinct. Would default to KPIs and process frameworks. Not equipped to be the cultural steward of a GDS FDE practice.
NOTES _____	

3 Signature FDE Questions

LEVEL 04 • MANAGER • THE RAREST ANSWERS COME FROM THE RAREST HIRES

1

THE MULTIPLIER TEST

"Tell me about an engineer you developed who you would now hire above yourself for a specific, named role. What did you do to get them there? And what did it cost you — personally or professionally — to develop them that well?"

WHAT YOU'RE LISTENING FOR

- Is the person real and named — or a hypothetical?
- Is there genuine pride in someone else's success — or is it self-congratulatory?
- Is the "cost" real — did developing them require something of the candidate personally?
- Do they frame being surpassed as success — or is there a defensive undertone?

GREEN SIGNALS

- ✓ Names a specific person without hesitation
- ✓ Describes the development work in detail — what was hard, what didn't work first
- ✓ The "cost" is real and acknowledged: time, ego, political capital
- ✓ Being surpassed is framed as the best outcome — and they mean it

RED SIGNALS

- ✗ Cannot name a specific person
- ✗ The "cost" is trivial or absent — "it just took time"
- ✗ The answer centres on the candidate's effort rather than the engineer's growth
- ✗ Being surpassed is answered as a hypothetical — they have not actually been surpassed

2

THE CULTURAL PRESERVATION AT SCALE TEST

"You lead a GDS FDE practice of 40 engineers. Three months in, the teams are performing — technically excellent, hitting milestones, clients are happy. But you notice something: they have stopped questioning whether they are building the right thing. They are executing brilliantly toward outcomes that nobody has recently verified are the right ones. Nobody has done anything wrong. How do you diagnose and address this before it becomes a problem?"

WHAT YOU'RE LISTENING FOR

- Do they recognise this as a cultural problem — not a process one?
- Is their intervention at the level of instinct and example — or policy and KPI?
- Do they understand why this happens — that comfort and momentum are the enemy of FDE culture?
- Is their diagnosis methodology specific and systemic?

GREEN SIGNALS

- ✓ Diagnoses it as a cultural drift, not a process failure
- ✓ Intervenes through example — does something visible that embodies the FDE instinct
- ✓ Uses a specific engagement as a diagnostic lens — not a survey or NPS metric
- ✓ Knows that "doing well" is the most dangerous time for cultural erosion

RED SIGNALS

- ✗ Introduces a new process or framework as the solution
- ✗ Does not see the danger because "the work is good and clients are happy"
- ✗ Solution is top-down: adds a "challenge the problem" step to the delivery checklist
- ✗ Has not thought about cultural preservation as a distinct leadership concern

3

THE CONSTRUCTIVE CHALLENGE TEST

"What is one thing about the way EY GDS FDE engagements are currently run that you would change on Day 1 of this role? And where do you think you might be wrong about that?"

WHAT YOU'RE LISTENING FOR

- Have they thought about the practice deeply enough to have a specific view?
- Is the challenge constructive — based on real observation — or a critique for its own sake?
- Can they challenge upward and hold the challenge simultaneously with genuine humility?
- Does "where I might be wrong" reveal actual intellectual honesty?

GREEN SIGNALS

- ✓ Specific, informed change — shows they have done their homework
- ✓ The change is grounded in a real problem they have observed, not a generic improvement
- ✓ "Where I might be wrong" is specific and shows they have genuinely considered it
- ✓ The humility is structural — they acknowledge what they do not yet know

RED SIGNALS

- ✗ "I'd learn before changing anything" — evasive non-answer
- ✗ The change is generic: "more collaboration," "clearer processes"
- ✗ "Where I might be wrong" is performative — a token caveat without real content
- ✗ Cannot identify a specific, substantive challenge — suggests insufficient preparation or thinking

Overall Assessment Summary

DIMENSION	SCORE (1-4)	WEIGHT
D1 · Technical Authority & Currency		×1.5
D2 · People Leadership & Talent Dev.		×1.5
D3 · Practice Building & Commercial		×1.0
D4 · High-Stakes Engagement Navigation		×1.0
D5 · FDE Culture & Multiplier Effect		×1.0
Signature Questions (avg)		×2.0
Weighted Average Score		Min 3.2 to Hire

HIRE / NO-HIRE SIGNALS

— Strong Hire Signals

- ✓ Names an engineer they'd hire above themselves
- ✓ Cultural drift answer is mechanism-based, not process-based
- ✓ Technically current enough to do a live architecture review
- ✓ Day-1 change is specific and shows real preparation

— Hard Pass Signals

- ✗ Cannot name anyone they developed who surpassed them
- ✗ Cultural preservation = new process or framework
- ✗ Last hands-on contribution was 18+ months ago
- ✗ Day-1 change is generic or evasive

INTERVIEWER DECISION

HIRE

HOLD

PASS

Agentic Systems *Engineering*

A specialist question bank for evaluating depth in agentic AI architecture. Use alongside the standard evaluation dimensions for L2–L4 candidates, or as a standalone technical screen for any level.

Questions are tagged with the minimum level at which each is appropriate.

01 • AGENT PATTERN SELECTION

02 • DEEP VS SHALLOW AGENTS

03 • HARNESS ENGINEERING

04 • SOLID IN AGENTIC DESIGN

05 • CONTEXT ENGINEERING

06 • DOMAIN FINE-TUNING PATTERNS

HOW TO USE THIS SECTION

L1 Candidates: Use 1–2 questions from each topic — awareness and recognition level only. You are testing whether they know a concept exists and can reason about it simply. Do not expect design depth.

L2–L3 Candidates: Use 2–3 questions per topic across all six areas. Expect design decisions with tradeoff reasoning. The best signal is when they draw on a specific production experience to answer.

L4 Candidates: Use all questions at full depth. Expect architectural opinions, challenge to standard patterns, and the ability to articulate when the standard approach is wrong. A strong L4 should disagree with at least one question premise.

01 Agent Pattern Selection — Simple to Multi-Agent

Can they select the right agent architecture for a given problem? Do they understand when a single well-prompted agent beats a five-agent orchestration? Do they know the failure modes of each pattern before implementing them?

L1

L2

L3

L4

"A client needs an AI system that: answers product questions, checks order status, processes refunds, and escalates to human agents when confident it cannot help. Walk me through the agent architecture you'd design. What pattern do you use and why?"

WHAT YOU'RE LISTENING FOR

- Do they reach for multi-agent when a router + tools suffices?
- Do they think about the human escalation path first or last?
- Do they ask about volume, latency requirements, and edge case frequency before designing?
- Do they name the pattern (orchestrator-worker, linear chain, router) or just describe it?

STRONG ANSWER

- ✓ Proposes a single orchestrator with 4 specialised tool-use agents
- ✓ Distinguishes between tool-use (stateless, fast) vs sub-agent (stateful, complex) for each task
- ✓ Designs the escalation path as a first-class agent capability — not an afterthought
- ✓ Asks about volume and SLA before committing to a pattern

WEAK ANSWER

- ✗ Immediately proposes 4 separate agents for 4 tasks without evaluating whether agents are needed at all
- ✗ Does not think about state management between intents
- ✗ Human escalation is "we'll add that later"
- ✗ Cannot name the pattern they're describing

L1 Calibration: Can they describe what an agent is and why you might use more than one? **L4 Push:** "When would you NOT use an agentic pattern for this problem at all? What would a non-agentic version look like and what would you trade off?"

L2

L3

L4

"Compare supervisor-agent, parallel-agent, and sequential-chain patterns. Give me a real problem that is best solved by each — and explain what makes each the wrong choice for the other two problems."

WHAT YOU'RE LISTENING FOR

- Does their example actually require the pattern — or could a simpler approach work?
- Do they understand that parallel agents require independent subtasks?
- Do they know that sequential chains accumulate latency at every step?
- Does the "wrong choice" explanation show real tradeoff thinking?

STRONG ANSWER

- ✓ **Supervisor:** Research task where sub-agents explore different domains and a supervisor synthesises — requires judgment on aggregation
- ✓ **Parallel:** Multi-source data ingestion where sources are independent and results merged — latency-optimised
- ✓ **Sequential:** Pipeline where each step enriches a document for the next — step N requires step N-1 output

WEAK ANSWER

- ✗ Uses vague examples that could apply to any pattern
- ✗ "Wrong choice" explanations are generic ("it would be slower")
- ✗ Confuses parallel execution with parallel agents
- ✗ Cannot name a failure mode specific to each pattern

- ✓ Can articulate the failure modes specific to each pattern

Key failure modes to probe: *Supervisor* — token cost explosion from full context passing. *Parallel* — race conditions on shared state. *Sequential* — error amplification where a bad step-2 output corrupts all downstream steps.

L2

L3

L4

"A colleague proposes a 12-agent swarm for a document classification task. You think it's over-engineered. Walk me through how you evaluate whether the complexity is justified — and what you would propose instead."

WHAT YOU'RE LISTENING FOR

- Do they have a principled framework for evaluating when multi-agent adds value vs. cost?
- Do they understand that complexity compounds — 12 agents means 12 failure surfaces?
- What does their "simpler alternative" look like? Is it actually simpler?

STRONG ANSWER

- ✓ Asks: What is the classification taxonomy? Are classes truly independent or overlapping?
- ✓ Proposes: single fine-tuned classifier or a 2-tier router + specialist for ambiguous cases
- ✓ Names the cost: 12x latency, 12x token cost, 12x failure surface, debugging complexity
- ✓ Has a clear "when would 12 agents be right?" — e.g. 12 genuinely distinct domains with no overlap

WEAK ANSWER

- ✗ Accepts the swarm approach as valid without questioning the premise
- ✗ Alternative is "fewer agents" without a principled reason for the number
- ✗ Does not raise the cost implications of the swarm approach
- ✗ Cannot articulate when the 12-agent approach WOULD be right

The FDE instinct here: *The best answer starts with "what problem are we actually solving?" not "how do we optimise the agent count?" Cargo cult agentic design is one of the most common failure modes in enterprise AI.*

L3

L4

"Design the complete agent architecture for the Meridian Bank fraud alert triage system — from alert ingestion through MRM-compliant decisioning to analyst workflow. Name every agent, its pattern type, scope boundary, and the failure mode you designed it to avoid."

WHAT YOU'RE LISTENING FOR

- Does each agent have a clearly bounded, single responsibility?
- Is the MRM compliance constraint reflected in the architecture — not added later?
- Does the failure mode reasoning show production experience or academic awareness?
- Is the analyst as a first-class participant in the design?

STRONG ANSWER

- ✓ Triage agent (router pattern) — avoids: routing the full 31K to analyst queue
- ✓ Resolution assist agent (tool-use pattern) — avoids: analyst starting from zero context
- ✓ Escalation agent (threshold + routing) — avoids: high-confidence fraud sitting in general queue
- ✓ Feedback loop agent (stateless extractor) — avoids: analyst decisions being lost to model team
- ✓ MRM compliance baked into every agent: HITL, hard thresholds, explainability chain

WEAK ANSWER

- ✗ Generic agent names without scope boundaries ("fraud agent", "analysis agent")
- ✗ MRM compliance appears as a separate step — "then we get approval"
- ✗ Failure modes are all technical (latency, cost) — no operational or regulatory failure modes
- ✗ Analyst appears at the end only — system is designed for efficiency, not human judgment

L4 extension: *"Which of these agents could be replaced with a well-designed rule-based system without any LLM — and would that be better?" Strong L4 candidates will know the answer is yes for at least one of them.*

02 Deep Agents vs Shallow Agents

A deep agent reasons across many steps, calls tools recursively, and spawns sub-agents. A shallow agent executes one or two steps with high reliability and low latency. Neither is better — knowing which to use is the skill. Candidates who always reach for deep agents have never debugged one in production.

L1

L2

L3

L4

"Define a deep agent and a shallow agent. Give me a real example of each from something you have built or studied. When is one clearly the wrong choice?"

WHAT YOU'RE LISTENING FOR

- Can they give a crisp, accurate definition of both?
- Are their examples real — or theoretical constructs?
- Do they understand depth as a risk multiplier, not just a capability?
- Can they name a scenario where a shallow agent beats a deep one?

STRONG ANSWER

- ✓ Deep: multi-step research agent that queries sources, synthesises, forms hypotheses, queries again
- ✓ Shallow: alert triage agent that applies a classification rule and routes — no iterative reasoning needed
- ✓ Deep is wrong: any latency-sensitive, high-volume task where each step compounds error rate
- ✓ Shallow is wrong: open-ended tasks requiring dynamic problem decomposition

WEAK ANSWER

- ✗ Frames deep = better, shallow = limited
- ✗ Examples are both variations of the same pattern
- ✗ Cannot describe a scenario where depth is the wrong choice
- ✗ Conflates "deep" with "multi-agent" — these are orthogonal dimensions

L2

L3

L4

"You have a task requiring 14 sequential tool calls to complete. How do you decide between a single deep agent doing all 14 steps versus a shallow orchestrator routing to 7 specialised 2-step agents? What are the decisive factors?"

WHAT YOU'RE LISTENING FOR

- Do they frame this as a tradeoff analysis — not a binary choice?
- Do they raise error compounding as a concern for deep agents?
- Do they understand that shallow agents are more testable and debuggable?
- Do they think about what happens when step 8 fails in each architecture?

STRONG ANSWER

- ✓ Decisive factors: Are the 14 steps truly sequential or can groups be parallelised? Does each step require the full context of all previous steps or just the immediately prior output?
- ✓ Deep agent wins if: reasoning must span all 14 steps, each decision informs the next
- ✓ Shallow wins if: steps cluster into independent phases; error recovery is per-cluster not full restart
- ✓ Names the key risk of deep: reasoning drift after step 7-8 as context accumulates

WEAK ANSWER

- ✗ Chooses based on engineering convenience rather than problem structure
- ✗ Does not raise context accumulation as a depth-specific risk
- ✗ Does not think about testability — how do you unit-test step 9 of a 14-step agent?
- ✗ "It depends" without specifying what it depends on

The key insight to probe for: Shallow agents are independently testable, cacheable, and observable. Deep agents are a single debuggable unit — which sounds simpler, but becomes increasingly opaque as depth increases. This is the

L2

L3

L4

"What are the failure modes that emerge specifically in deep agent architectures that are not present or are much rarer in shallow agents? For each failure mode, what is the detection signal and what is the design mitigation?"

WHAT YOU'RE LISTENING FOR

- Do they name failure modes from production experience — not just theory?
- Is their detection mechanism specific and implementable?
- Does their mitigation actually address the root cause?

STRONG FAILURE MODES NAMED

- ✓ Reasoning drift: Agent's goal at step 12 has imperceptibly shifted from step 1. Detection: compare step-1 intent with step-12 plan. Mitigation: inject original objective reminder every N steps
- ✓ Runaway loops: Agent keeps calling the same tool expecting a different result. Detection: tool-call dedup + loop counter. Mitigation: hard call limit per tool per session
- ✓ Context poisoning: A bad tool output in step 4 corrupts all downstream reasoning. Detection: output validation at each step. Mitigation: tool output schema enforcement
- ✓ Unbounded execution: Agent cannot determine when "done" — keeps extending the plan. Mitigation: explicit termination condition in system prompt + step budget

WEAK ANSWER

- ✗ Names only generic failures (hallucination, latency) that apply to all agents
- ✗ Detection mechanisms are "check the logs" without specificity
- ✗ Mitigations are all post-hoc — no design-time mitigations proposed
- ✗ Cannot name a failure mode they have personally encountered

03 Harness Engineering

A harness is the infrastructure that wraps an agentic system to make it testable, observable, and safe to run. Eval harnesses, test harnesses, monitoring harnesses, and safety harnesses are all distinct concerns. The best FDE engineers build the harness before the agent. Most engineers build it after the incident.

L1

L2

L3

L4

"What is a harness in the context of an agentic AI system? Describe the four types of harness you would build around a production fraud triage agent — and why each is a separate concern."

WHAT YOU'RE LISTENING FOR

- Do they know that harnesses are a category — not a single thing?
- Can they articulate why each type is separately implemented?
- L1: awareness that harnesses exist. L3+: can design each from scratch

FOUR HARNESSES

- ✓ Test Deterministic, harness: repeatable — verifies agent produces correct output on known inputs. Runs offline. Input/output fixtures. Regression suite.
- ✓ Eval Probabilistic quality harness: measurement — precision, recall, hallucination rate, latency percentiles. Runs against golden datasets. Informs release decisions.
- ✓ Monitoring Production harness: observability — live metrics, drift detection, alert volume tracking, anomaly detection on agent behaviour.
- ✓ Safety Guardrail enforcement harness: — hard thresholds, HITL triggers, PII redaction, output validation, circuit breakers.

WEAK ANSWER

- ✗ Describes "harness" as a single thing — usually just a test suite
- ✗ Monitoring and eval are treated as the same concern
- ✗ Safety harness is not mentioned or is equated with "adding guardrails later"
- ✗ Cannot explain why separating them matters operationally

L2

L3

L4

"Walk me through how you build an eval harness for a new fraud triage agent from scratch — before the agent exists. What does the harness contain, how do you populate the golden dataset, and what is the release gate definition?"

WHAT YOU'RE LISTENING FOR

- Do they build the harness before the agent — or as an afterthought?
- How do they construct a golden dataset when no agent output exists yet?

STRONG ANSWER

- ✓ Harness built before agent — defines what success looks like before implementation
- ✓ Golden dataset sourced from historical analyst decisions (labelled ground truth)

WEAK ANSWER

- ✗ Builds the harness after building the agent to "check it works"
- ✗ Golden dataset is "sample data we collected" without labelling methodology
- ✗ Release gate is "accuracy > 80%" without specifying what

→ Is the release gate specific and measurable — or vague? → Do they understand that the harness is the specification of what "working" means?	✓ Metrics: precision/recall per alert class, P95 latency, hallucination rate on risk narrative ✓ Release gate: precision ≥ 0.94 on P1 alerts, false negative rate $\leq 0.02\%$, latency P95 $< 800\text{ms}$ ✓ Human review of a stratified 5% sample before production	- accuracy means ✗ No human review step in the release process
--	---	---

The harness-first principle: *If you build the agent before the eval harness, you are writing code for a specification that does not yet exist. The harness IS the specification — it makes "done" measurable before the first line of agent code is written.*

L3 L4 "You inherit a production agentic system that has been running for 4 months with no harness infrastructure — no eval, no monitoring, no safety layer. It is handling real transactions. How do you retrofit harness infrastructure without breaking the running system?"		
WHAT YOU'RE LISTENING FOR → Do they sequence correctly — safety harness first, then monitoring, then eval? → Do they understand the risk of retrofitting vs. greenfield? → How do they build a golden dataset from a system with no historical labels? → Do they use shadow deployment / parallel run as an interim safety mechanism?	STRONG ANSWER ✓ Safety harness first: add hard output validators and circuit breakers before anything else — these can be added as a pass-through layer without touching the agent ✓ Add structured logging to capture inputs/outputs of every agent call — this is the raw material for the eval and monitoring harnesses ✓ Build monitoring from the logs: alert on statistical anomalies in output distribution before you have golden labels ✓ Build golden dataset retrospectively using human review of a stratified sample of captured outputs ✓ Eval harness last — once you have labelled data and baseline metrics	WEAK ANSWER ✗ Starts by adding eval harness — without safety coverage, this risks finding problems you cannot safely act on ✗ Plans a "harness sprint" that requires taking the system down ✗ Does not consider the risk of adding instrumentation to a live system ✗ No plan for building a golden dataset retrospectively

04 SOLID Principles in Agentic Engineering

The SOLID principles apply to agentic systems with precise and practical implications. Most candidates can recite SOLID from memory. The test is whether they can apply each principle to a real agentic design decision — and recognise when violating one is causing a live problem.

L2

L3

L4

"Apply the Single Responsibility Principle to agent design. Show me a concrete example of an agent that violates SRP and one that respects it — using the fraud triage system as context. What breaks in production when SRP is violated?"

WHAT YOU'RE LISTENING FOR

- Is their SRP violation genuinely multi-responsibility — or just "too large"?
- Can they articulate the production consequence of the violation?
- Does their compliant version maintain testability as a natural property?

STRONG ANSWER

- ✓ Violation: A single "fraud ops agent" that triages alerts, retrieves case history, generates risk narrative, routes to analyst, AND packages model feedback — six responsibilities, one agent
- ✓ Consequence: Changing the routing logic requires re-testing the narrative generation. A prompt change for narratives breaks the triage logic. Impossible to unit-test triage in isolation.
- ✓ Compliant: Four agents — triage agent, case enrichment agent, narrative generation agent, feedback packaging agent — each independently testable and deployable

WEAK ANSWER

- ✗ SRP violation is described as "too big a prompt" — confuses size with responsibility
- ✗ Cannot name a specific production consequence of the violation
- ✗ Compliant version is "a smaller agent" without naming the separated responsibilities
- ✗ Testability is not mentioned as a benefit of SRP compliance

L2

L3

L4

"How does the Dependency Inversion Principle apply to tool architecture in an agentic system? Give me a concrete example from a banking context where violating DIP caused or would cause a real problem."

WHAT YOU'RE LISTENING FOR

- Do they understand DIP as "depend on abstractions, not implementations"?

STRONG ANSWER

- ✓ Violation: Orchestrator `salesforce_crm_api_v3.get_customer(id)` directly — it depends on a specific implementation

→ Is their banking example concrete — or a restatement of the principle? → Do they connect DIP to the real operational risk of vendor lock-in or API versioning?	✓ Problem: When the bank upgrades CRM from Salesforce to SAP, every agent that calls the old API must be rewritten and re-tested ✓ DIP compliant: Agent <code>get_customer_profile(id: str) → CustomerProfile</code> — an abstraction. The CRM implementation behind it can change without touching the agent. ✓ Extends to LLM providers: <code>generate_text(prompt)</code> — <code>anthropic_claude_3_opus(prompt)</code> not on agent depends on
--	--

L2 L3 L4

"How does the Open/Closed Principle apply when you need to add new tool capabilities to an agentic system that is already running in production? What is the architectural mechanism that allows extension without modification?"

WHAT YOU'RE LISTENING FOR

- Do they understand OCP as "open for extension, closed for modification"?
- Can they describe a tool registry pattern or plugin architecture?
- Do they connect this to MCP (Model Context Protocol) or equivalent?

STRONG ANSWER

- ✓ Mechanism: A tool registry that agents query at runtime — tools are registered, not hardcoded into agent logic
- ✓ Adding a new tool = registering it with the registry and defining its schema. The agent discovers it via the registry without modification.
- ✓ MCP is the natural embodiment of OCP for agentic systems — server-defined tools consumed by agents without agent code changes
- ✓ Closed for modification means: adding a new tool should not require changing the agent's system prompt, code, or deployment

WEAK ANSWER

- ✗ "We'd update the system prompt with the new tool description" — this IS modification
- ✗ Does not describe a registry or discovery mechanism
- ✗ No mention of how tool addition would be tested without touching existing agent

The Liskov and Interface Segregation extensions (L4): LSP — can you swap Claude for GPT-4 in your agent and get equivalent behaviour? If not, your agent is coupled to a specific model's idiosyncrasies. ISP — are your tools focused and single-purpose, or do they return a monolithic data blob that agents must parse?

Context management is the most underestimated engineering challenge in production agentic systems. Context rot, context overload, context drift, and context leakage are all distinct failure modes with distinct signatures. The best engineers design context management as a first-class architectural concern — not a tuning parameter.

L1

L2

L3

L4

"What is context rot in an agentic system? What does it look like in production — describe the observable symptom a non-technical stakeholder would notice before any engineer looks at a log?"

WHAT YOU'RE LISTENING FOR

- Do they distinguish context rot from context overload?
- Can they describe the symptom from the user's perspective — not the engineer's?
- Is their understanding from production observation or academic reading?

STRONG ANSWER

- ✓ **Definition:** Context rot is the degradation of reasoning quality as stale, irrelevant, or contradictory information accumulates in the active context window over a long session
- ✓ **Observable symptom:** The agent begins giving answers that contradict decisions it made earlier in the same session. A fraud analyst asks the same type of question they asked in step 3 and gets a different answer in step 15. The agent seems to "forget" the constraints it was given at the start.
- ✓ **Distinct from context overload** (which is hitting the token limit) — rot can happen well within the window limit

WEAK ANSWER

- ✗ Conflates context rot with context overload (running out of tokens)
- ✗ Symptom description is engineer-level ("I see inconsistent outputs in the logs") not user-level
- ✗ Cannot distinguish rot as a quality failure from a capacity failure

L2

L3

L4

"Walk me through your strategy for managing context in a fraud operations agent that runs for 2+ hours per analyst session, processing dozens of alerts. What techniques do you use and at what point in the session does each kick in?"

WHAT YOU'RE LISTENING FOR

- Do they have a layered strategy — not a single technique?
- Is their approach designed for graceful degradation — not just handling the common case?

STRONG TECHNIQUES

- ✓ **Progressive summarisation:** Every N alerts, compress the session context into a structured summary. Inject the summary,

WEAK ANSWER

- ✗ "Use a larger context window model" — avoids the problem rather than solving it
- ✗ Single technique applied uniformly — no layered strategy

→ Do they separate session context from historical context from system context?	drop the raw history. ✓ Selective retention: Keep only decision-relevant context — the current alert, the last decision, the session-level constraints. Discard intermediate reasoning. ✓ External state store: Move session state (analyst decisions, running fraud totals, flagged patterns) out of the context window into a structured store. Retrieve on demand. ✓ Context horizon segmentation: Immediate context (current alert), session context (today's decisions), historical context (retrieved from vector store). Each managed separately. ✓ Soft and hard context budgets: At 60% of context window — summarise. At 80% — compress aggressively. At 90% — start new session with carried-over summary.	✗ Summarisation is their only tool and it is not scheduled — triggered only when they remember ✗ Does not separate the three context horizons
--	--	---

L3

L4

"Name and describe the five most dangerous context failure modes in production agentic systems. For each: the observable symptom, the detection method, and the architectural mitigation."

WHAT YOU'RE LISTENING FOR

- Do they name all five or can they get to at least three with real depth?
- Are their detection methods instrumentable — not just "check the logs"?
- Are mitigations design-time decisions — not post-hoc patches?

FIVE FAILURE MODES

- ✓ Context Overload: Exceeds window limit
→ agent truncates early instructions → injects counter-instructions into user turn. Detect: token counter approaching 85% threshold. Mitigate: progressive summarisation + hard session segmentation.
- ✓ Context Rot: Stale or irrelevant content degrades reasoning quality. Detect: semantic similarity between current output and expected output from

WEAK ANSWER

- ✗ Names fewer than three distinct failure modes
- ✗ Conflates overload and rot
- ✗ Context leakage and context poisoning are absent — these are the least-discussed but most dangerous
- ✗ Mitigations are all reactive ("restart the session") rather than preventive design choices

	same prompt at session start. Mitigate: selective retention + scheduled compression. ✓ Context A malformed tool Poisoning: output or adversarial input contaminates the context. Detect: output schema validation on every tool call. Mitigate: sandboxed tool execution + output sanitisation layer. ✓ Context System prompt Drift: constraints are overridden or diluted over a long session. Detect: periodically re-inject system prompt and compare agent behaviour. Mitigate: system prompt pinning + constraint re-assertion at N-step intervals. ✓ Context Sensitive information Leakage: from one user session bleeds into another via a shared context store. Detect: session isolation test — confirm session A cannot see session B outputs. Mitigate: strict session-scoped context namespacing + TTL-based context expiry.	
--	---	--

06 Domain-Specific Fine-Tuning Patterns

Fine-tuning is the most over-applied and least-understood technique in enterprise AI. Most clients who ask for a fine-tuned model actually need better RAG, better prompts, or a different task framing. The FDE engineer's job is to know the difference — and when fine-tuning is genuinely the right answer, to execute it with rigour.

L1

L2

L3

L4

"A Meridian Bank stakeholder says: 'Our fraud analysts use very specific language and frameworks when writing case summaries — we need the AI to write exactly like them. Can we fine-tune the model?' Walk me through your decision process for whether fine-tuning is the right approach."

WHAT YOU'RE LISTENING FOR

- Do they evaluate alternatives before committing to fine-tuning?
- Is their decision framework structured — not just instinct?
- Do they understand that "write like us" is a style/format problem — often solvable with prompting?
- Do they ask about data availability before recommending fine-tuning?

DECISION FRAMEWORK

- ✓ Step 1: Can prompt engineering + few-shot examples solve it? "Write like these 5 example summaries" — if 5 examples achieve 80% of the target style, fine-tuning ROI is low.
- ✓ Step 2: Is this a style/format problem or a knowledge/domain problem? Style → prompts first. Domain knowledge (acronyms, entities, proprietary frameworks) → RAG or fine-tune.
- ✓ Step 3: Data availability — do we have 500+ high-quality examples? If not, fine-tuning will overfit. What is the annotation cost?
- ✓ Step 4: Latency and cost constraints — fine-tuned smaller model may be faster and cheaper than GPT-4 with a long few-shot prompt.
- ✓ Verdict: For "write like our analysts" — start with few-shot prompting. Fine-tune only if prompting plateaus and you have 1,000+ quality examples.

WEAK ANSWER

- ✗ Immediately agrees fine-tuning is the right approach without exploring alternatives
- ✗ Does not ask about data availability
- ✗ Does not distinguish style fine-tuning from domain knowledge fine-tuning
- ✗ Cannot describe what "fine-tuning plateau" means — when you'd know prompting has been exhausted

L2

L3

L4

"Design a fine-tuning data pipeline for a model that needs to generate fraud investigation narratives in Meridian Bank's house style. Include: data sourcing, quality validation, augmentation strategy, train/eval split design, and the trigger conditions for retraining."

WHAT YOU'RE LISTENING FOR

STRONG PIPELINE DESIGN

WEAK ANSWER

→ Is their data sourcing strategy realistic — where does the data actually come from? → Is quality validation automated, human-reviewed, or both? → Do they account for class imbalance in fraud data? → Is their retraining trigger based on performance metrics — not a calendar?	✓ Source: Historical analyst-written case summaries (3+ years), labelled with alert type, fraud category, and analyst seniority. Filter: only summaries with confirmed outcome labels. ✓ Quality validation: Automated: length filter, language detection, PII scan. Human: 10% sample reviewed by senior analyst for accuracy and style conformance. ✓ Augmentation: Paraphrase low-frequency fraud categories using base model + prompt to address class imbalance. Do not augment majority class. ✓ Split: Time-based — last 6 months as held-out eval (represents current language patterns). Not random split — random split leaks future patterns into training. ✓ Retraining trigger: Style conformance score drops below 0.85 on weekly eval sample, OR a new fraud pattern emerges with >200 labelled examples, OR language model update requires re-alignment.	✗ Data source is vague — "we collect examples from analysts" without describing the labelling process ✗ Random train/eval split — does not recognise temporal leakage risk in fraud data ✗ No augmentation strategy for rare fraud categories ✗ Retraining is time-based ("monthly") rather than performance-triggered
--	---	---

L2 L3 L4

"Your fine-tuned fraud narrative model performs excellently on the evaluation set — 94% style conformance, low hallucination rate. Two weeks into production, it begins catastrophically failing on a category of alerts it handled well before fine-tuning. What happened and how do you address it?"

WHAT YOU'RE LISTENING FOR → Do they immediately identify catastrophic forgetting as the mechanism? → Do they know the standard mitigations — rehearsal data, LoRA, elastic weight consolidation?	STRONG ANSWER ✓ Diagnosis: Catastrophic forgetting — the fine-tuning process overwrote the base model's general capability for the alert category that was under-	WEAK ANSWER ✗ Does not name catastrophic forgetting ✗ Attributes failure to "insufficient training data" without explaining the mechanism ✗ Mitigation is "collect more data" — misses the architectural solutions (LoRA, rehearsal)
--	--	---

→ Is their root cause analysis specific — not just "it forgot"?	represented in the fine-tuning dataset ✓ Root cause: The fine-tuning data had very few examples of this alert category, so the model's gradient updates for the dominant categories over-wrote the weights governing rare-category behaviour ✓ Mitigation options: (1) Rehearsal data: include representative examples from the base model's original capability distribution in fine-tuning. (2) LoRA: fine-tune only adapter layers, preserve base model weights — standard mitigation for this exact scenario. (3) Re-examine data balance: augment the under-represented category before re-tuning. ✓ Immediate action: roll back to the base model + few-shot prompt for the affected category while the fix is prepared	✗ No immediate rollback or containment plan — plans to fix before reverting
--	---	--

L3/L4 extension: "How does LoRA specifically protect against catastrophic forgetting compared to full fine-tuning — and what does this imply about where you should always use LoRA for domain adaptation in production systems?" Strong answer: LoRA adds low-rank adapter matrices without modifying base weights — the base model's general capabilities are frozen. This is the primary reason LoRA is the default choice for production domain adaptation.

L3

L4

"When is fine-tuning strictly better than RAG for domain adaptation — and when is it strictly worse? Give me a decision matrix with four real variables that determine the choice."

WHAT YOU'RE LISTENING FOR	FOUR DECISION VARIABLES	WEAK ANSWER
→ Is their decision matrix genuinely operational — or academic? → Do they understand that these techniques are often complementary rather than competitive? → Do they account for maintenance burden — fine-tuning requires retraining; RAG requires corpus management?	✓ Knowledge volatility: If domain knowledge changes frequently (market data, policy updates) → RAG wins. Knowledge is retrieved fresh. Fine-tuning goes stale immediately. ✓ Task type: If the task is style/behaviour/reasoning pattern adaptation (write like our analysts, follow our process) → Fine-tuning wins. RAG cannot change how the model	✗ Frames RAG vs fine-tuning as binary — does not consider the complementary pattern ✗ Decision variables are vague ("complexity of the domain", "size of the dataset") ✗ Does not mention knowledge volatility — the most operationally important variable ✗ No mention of maintenance burden for either approach over time

	reasons, only what it knows. ✓ Latency If sub-100ms requirements: response is required → Fine-tune a smaller model. RAG adds retrieval latency at every call. ✓ Data If you have <500 availability: high-quality labelled examples → RAG wins. Fine-tuning with insufficient data will overfit. RAG works with even a few dozen high-quality documents. ✓ Best practice: Fine-tune for behaviour/style + RAG for knowledge = the production pattern most high-performing enterprise AI systems use.	
--	---	--

EY GDS · Forward Deployment Engineering

INTERVIEW EVALUATION GUIDELINES · CONFIDENTIAL · INTERNAL USE ONLY
 OUTCOME MERCENARIES WITH TECHNICAL DEPTH